

```

# Step 1: Upload CSV in Colab
from google.colab import files
import pandas as pd

uploaded = files.upload() # Select your kidney_disease.csv

# Load dataset
df = pd.read_csv("kidney_disease.csv")
print("Shape:", df.shape)
df.head()

<IPython.core.display.HTML object>

Saving kidney_disease.csv to kidney_disease (1).csv
Shape: (400, 26)

{"type": "dataframe", "variable_name": "df"}

import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler

# Drop ID column if it exists
if 'id' in df.columns:
    df = df.drop('id', axis=1)

# Handle missing values (fill numeric with median, categorical with mode)
for col in df.columns:
    if df[col].dtype == 'object':
        df[col].fillna(df[col].mode()[0], inplace=True)
    else:
        df[col].fillna(df[col].median(), inplace=True)

# Encode categorical columns
le = LabelEncoder()
for col in df.columns:
    if df[col].dtype == 'object':
        df[col] = le.fit_transform(df[col])

# Define X and y
X = df.drop('classification', axis=1) # assuming 'classification' is target
y = df['classification']

# Scale features
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,

```

```
test_size=0.2, random_state=42, stratify=y)

print("Training shape:", X_train.shape)
print("Testing shape:", X_test.shape)

Training shape: (320, 24)
Testing shape: (80, 24)

/tmp/ipython-input-214255797.py:14: FutureWarning: A value is trying
to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.
```

```
df[col].fillna(df[col].median(), inplace=True)
/tmp/ipython-input-214255797.py:12: FutureWarning: A value is trying
to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.
```

```
df[col].fillna(df[col].mode()[0], inplace=True)

from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, classification_report

# Train model
model = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
random_state=42)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Evaluate
acc = accuracy_score(y_test, y_pred)
```

```

print("XGBoost Accuracy:", acc)
print(classification_report(y_test, y_pred))

/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183:
UserWarning: [14:05:54] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)

□ XGBoost Accuracy: 0.975
      precision    recall   f1-score   support
          0       0.96     1.00     0.98      50
          2       1.00     0.93     0.97      30

      accuracy           0.97      80
      macro avg       0.98     0.97     0.97      80
weighted avg       0.98     0.97     0.97      80

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Train
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)

# Predict
y_pred_rf = rf.predict(X_test)

# Evaluate
acc_rf = accuracy_score(y_test, y_pred_rf)
print("Random Forest Accuracy:", acc_rf)
print(classification_report(y_test, y_pred_rf))

□ Random Forest Accuracy: 0.9875
      precision    recall   f1-score   support
          0       0.98     1.00     0.99      50
          2       1.00     0.97     0.98      30

      accuracy           0.99      80
      macro avg       0.99     0.98     0.99      80
weighted avg       0.99     0.99     0.99      80

from lightgbm import LGBMClassifier
from sklearn.metrics import accuracy_score, classification_report

# Train with regularization + constraints
lgb = LGBMClassifier(

```


	macro avg	0.99	0.98	0.99	80
	weighted avg	0.99	0.99	0.99	80

```

/usr/local/lib/python3.12/dist-packages/sklearn/utils/
validation.py:2739: UserWarning: X does not have valid feature names,
but LGBMClassifier was fitted with feature names
  warnings.warn(
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import numpy as np # Import numpy

# 1. Class distribution
sns.countplot(x='classification', data=df)
plt.title("Class Distribution")
plt.show()

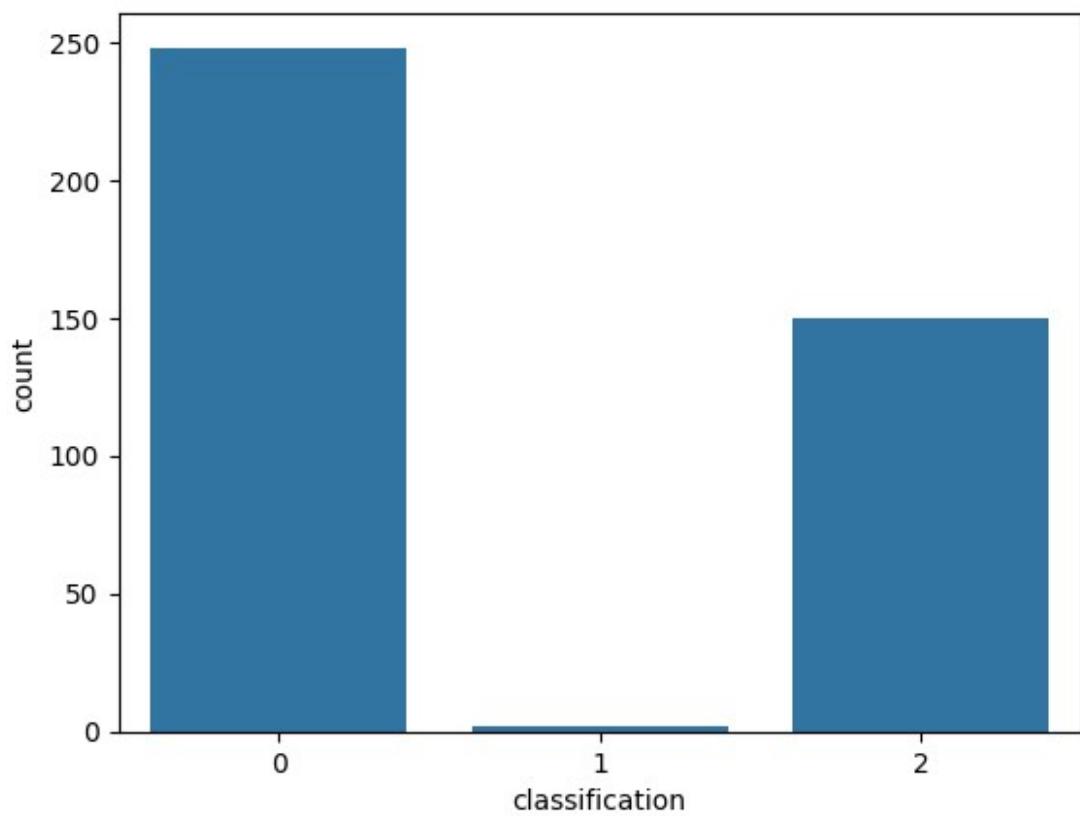
# 2. Correlation heatmap
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), cmap="coolwarm", annot=False)
plt.title("Feature Correlation Heatmap")
plt.show()

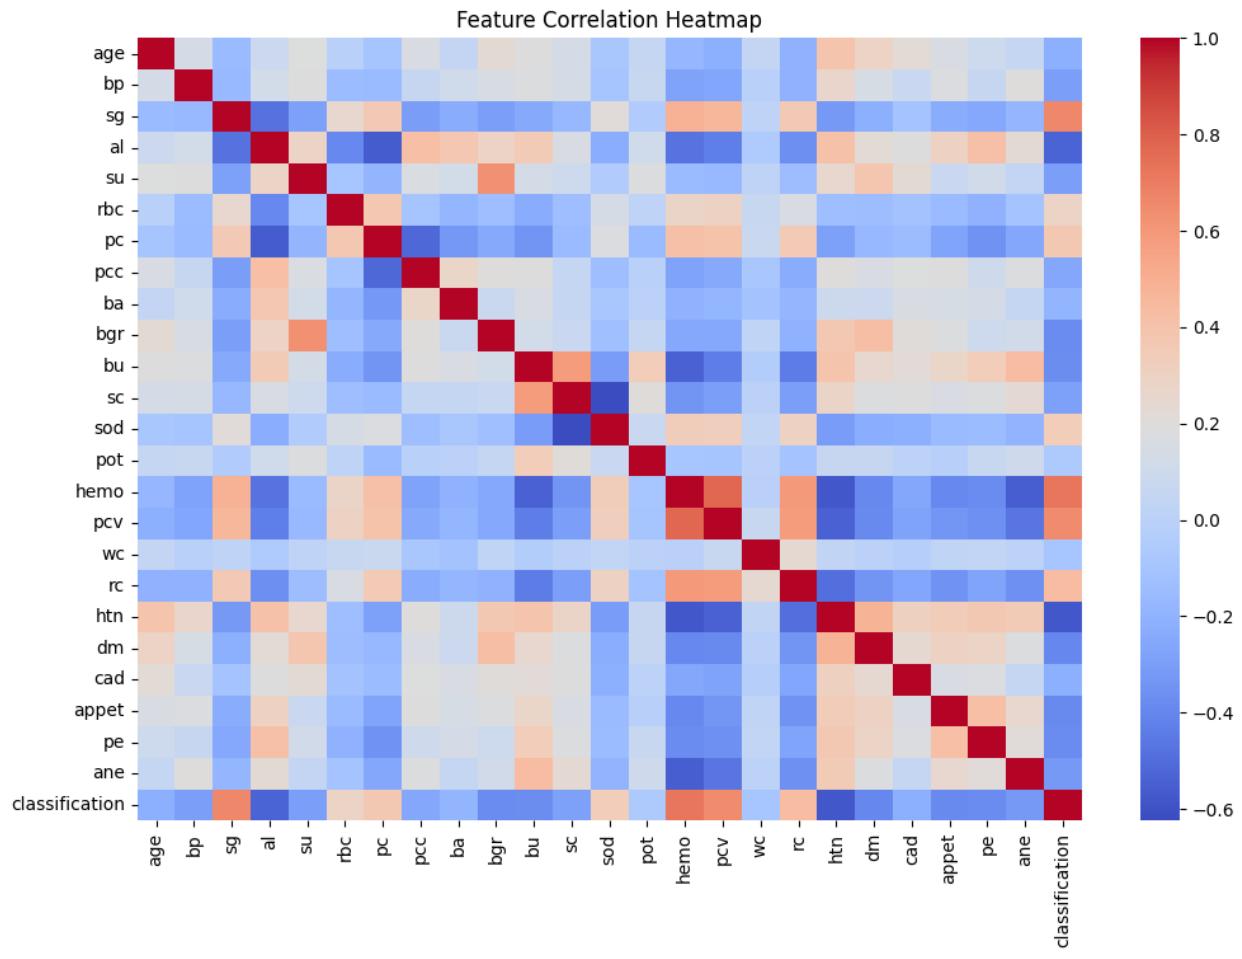
# 3. Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
# Use unique values from y_test for display labels
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=np.unique(y_test))
disp.plot(cmap="Blues")
plt.title("Confusion Matrix")
plt.show()

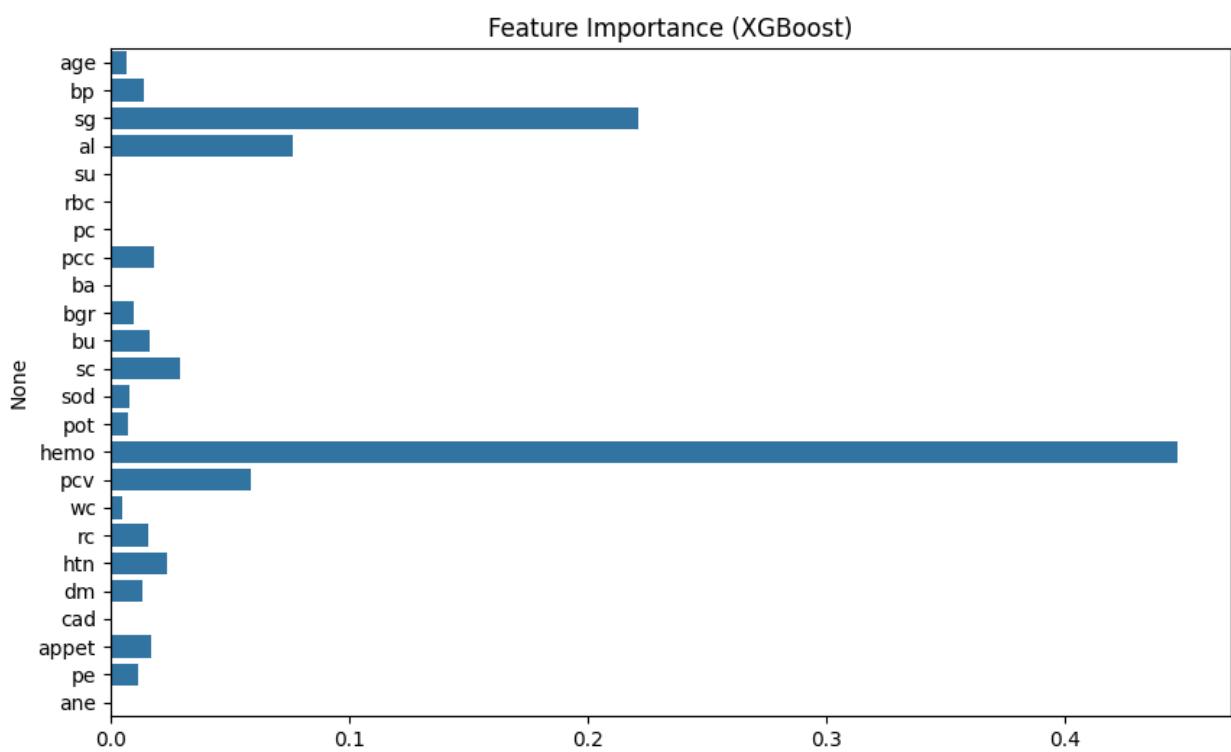
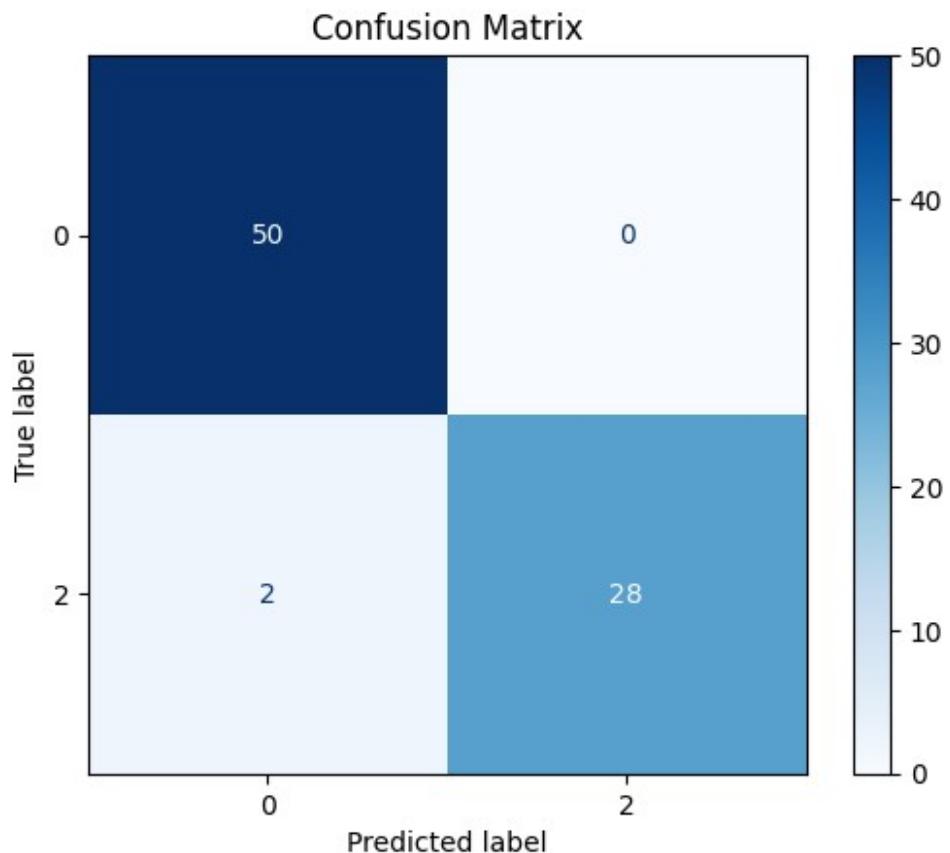
# 4. Feature Importance
plt.figure(figsize=(10,6))
sns.barplot(x=model.feature_importances_, y=df.drop('classification',
axis=1).columns)
plt.title("Feature Importance (XGBoost)")
plt.show()

```

Class Distribution







```

from google.colab import files
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from xgboost import XGBClassifier

try:
    df = pd.read_csv("kidney_disease.csv")
    print("Shape:", df.shape)
except FileNotFoundError:
    print("kidney_disease.csv not found. Please upload the file using the 'files.upload()' cell.")
    # Provide a placeholder DataFrame or exit if file not found
    df = pd.DataFrame() # Placeholder

if not df.empty:
    # Drop ID column if it exists
    if 'id' in df.columns:
        df = df.drop('id', axis=1)

        if df[col].dtype == 'object':
            df[col] = df[col].fillna(df[col].mode()[0])
        else:
            df[col] = df[col].fillna(df[col].median())

    le = LabelEncoder()
    for col in df.columns:
        if df[col].dtype == 'object':
            df[col] = le.fit_transform(df[col])

    X = df.drop('classification', axis=1) # assuming 'classification' is target
    y = df['classification']

    # Scale features
    scaler = StandardScaler()
    X = scaler.fit_transform(X)

    # Train-test split
    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42, stratify=y)

    print("Training shape:", X_train.shape)
    print("Testing shape:", X_test.shape)

    model = XGBClassifier(use_label_encoder=False,

```

```

eval_metric='logloss', random_state=42)
    model.fit(X_train, y_train)

Shape: (400, 26)
Training shape: (320, 24)
Testing shape: (80, 24)

/usr/local/lib/python3.12/dist-packages/xgboost/training.py:183:
UserWarning: [14:18:01] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

    bst.update(dtrain, iteration=i, fobj=obj)

# □ Histogram of Age
plt.figure(figsize=(8,5))
sns.histplot(data=df, x='age', bins=20, kde=True,
hue='classification')
plt.title("Age Distribution by CKD Status")
plt.show()

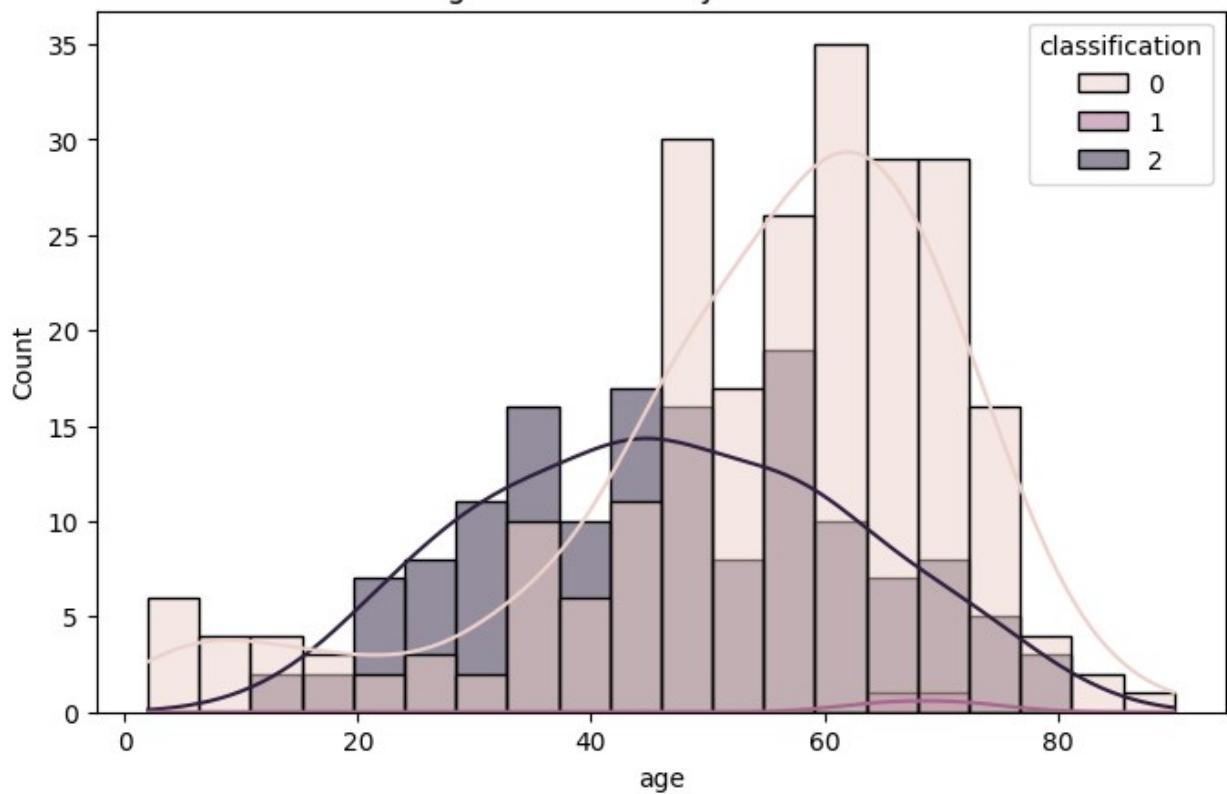
# □ Histogram of Blood Pressure
plt.figure(figsize=(8,5))
sns.histplot(data=df, x='bp', bins=20, kde=True, hue='classification')
plt.title("Blood Pressure Distribution by CKD Status")
plt.show()

# □ Scatter Plot: Age vs Blood Pressure
plt.figure(figsize=(8,5))
sns.scatterplot(x=df['age'], y=df['bp'], hue=df['classification'],
alpha=0.7)
plt.title("Age vs Blood Pressure (Colored by CKD Status)")
plt.show()

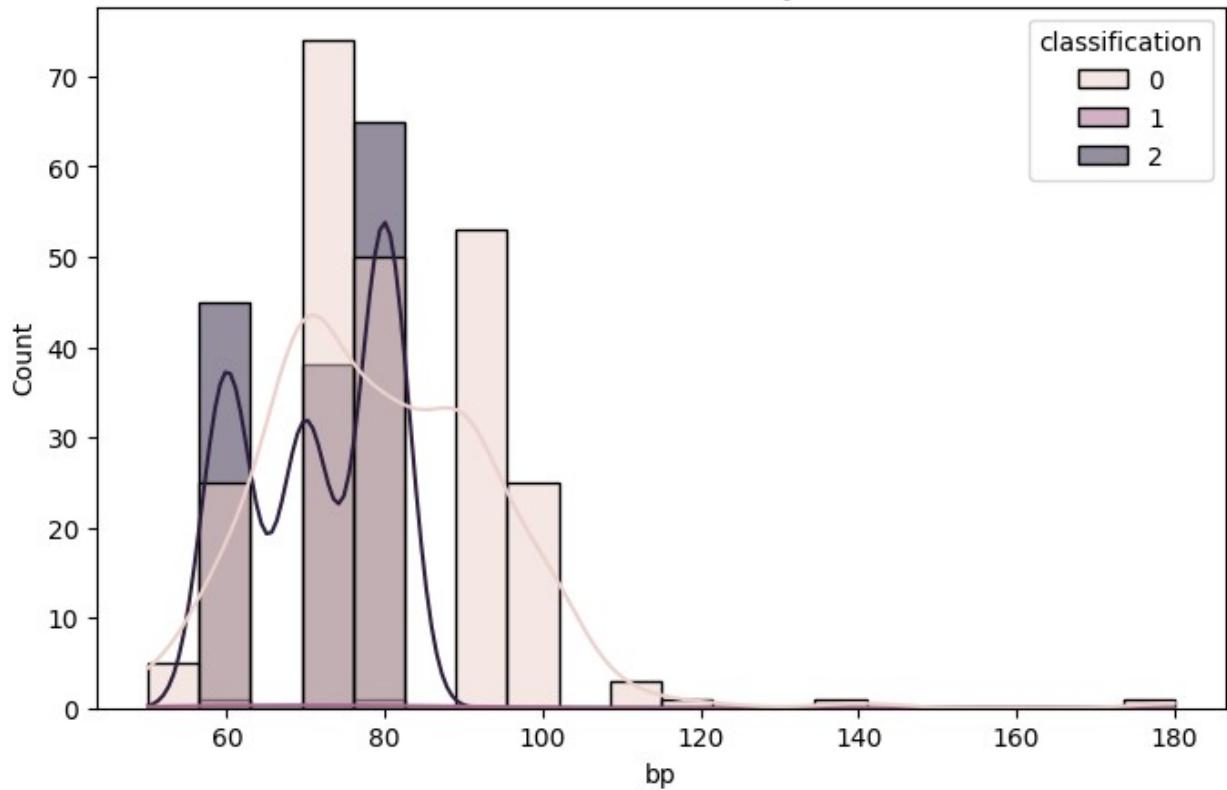
# □ Pairplot for selected features
selected_features = ['age', 'bp', 'sg', 'al', 'su', 'classification']
sns.pairplot(df[selected_features], hue='classification')
plt.suptitle("Pairplot of Selected Features", y=1.02)
plt.show()

```

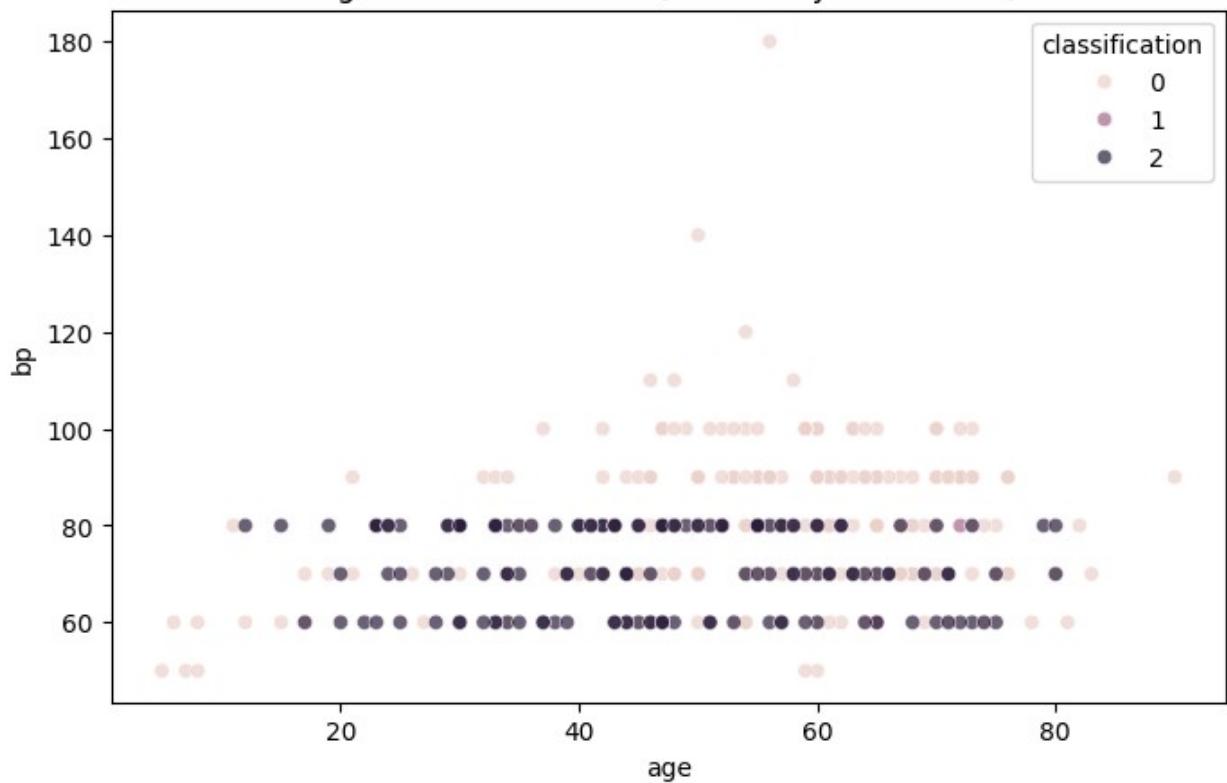
Age Distribution by CKD Status



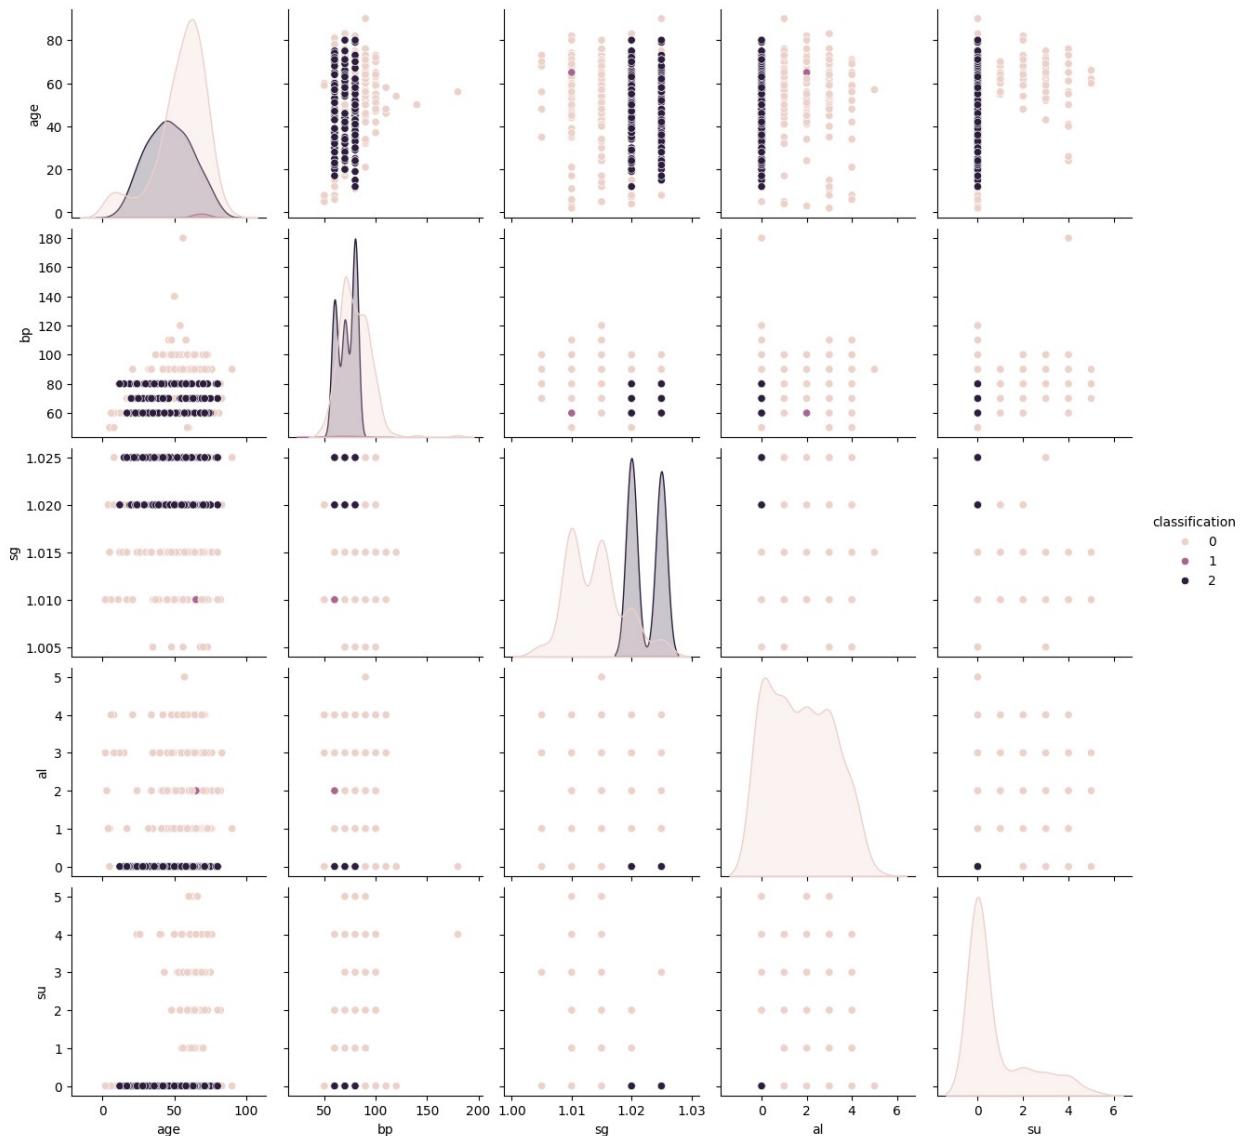
Blood Pressure Distribution by CKD Status



Age vs Blood Pressure (Colored by CKD Status)



Pairplot of Selected Features



```
import shap
explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values, X_test)
```

